

# SPEC-1—Dr. Patel 4th-Trimester Tutor Chatbot

## Background

**Problem.** New parents often face urgent, anxiety-provoking questions at 1-3 AM about a newborn's feeding, sleep, temperature, breathing, jaundice, and postpartum recovery. Dr. Sonal Patel (pediatrician/neonatologist; founder of NayaCare) focuses on comprehensive "4th trimester" care, but cannot be on-call 24/7 for non-emergent education and guidance.

**Vision.** A friendly, compassionate chatbot embedded on NayaCare's homepage that serves as an educational tutor and resource hub for postpartum families (0-12+ weeks after delivery). It should:

- Offer clear, empathetic, evidence-informed education for both infant and birthing parent topics.
- Triage concerns to **educational guidance vs. seek-care-now** messaging, without diagnosing or replacing medical care.
- Surface only **physician-vetted resources** (Dr. Patel's curated library) by default; optionally cite reputable clinical guidelines for generic, low-risk topics.
- Generate take-home **PDF summaries** and handouts personalized to the conversation context.
- Maintain **HIPAA compliance** (data minimization, secure storage/transit, auditability, BAAs with vendors), and default to *no PHI retention* beyond what's essential for safe operation.

**Primary users.** New parents/caregivers of newborns during the 4th trimester; secondary users include Dr. Patel and staff (content curation, analytics, quality review), and site visitors seeking general education.

**Content sources.** Dr. Patel's authored articles, clinic handouts, vetted PDFs, and selected guidelines. The bot prioritizes Dr. Patel's voice and materials; external sources are used conservatively and transparently.

**Non-goals (initial MVP).** No diagnosis, prescription, or treatment orders; no integration to EMR/PHI systems; no free-text intake of sensitive PHI unless strictly necessary for triage wording; no multilingual support beyond English (unless specified later).

**Success (MVP).** 24/7 availability; <60s time-to-useful-answer; high parent satisfaction; safe escalation wording for red-flag symptoms; minimal hallucinations; PDFs with correct, vetted content; zero known HIPAA incidents.

## Requirements

### Must Have (M)

- **On-site chat widget (web)** embedded on [nayacare.org](http://nayacare.org) homepage; mobile-first; WCAG 2.2 AA; clear emergency disclaimer & consent gate before first use.
- **Acceptance:** Loads in <2s on 4G; keyboard + screen-reader accessible; privacy notice shown before first message.
- **Scope guardrails:** education + risk-screening only; no diagnosis, no prescriptions, no dosing unless from Dr. Patel-approved handouts; always include "does not replace care" footer.
- **Acceptance:** 100% conversations include scope footer; red-flag flows never output a diagnosis.

- **Triage safety net** with newborn & maternal red-flags and automatic escalation messaging: e.g., infant fever  $\geq 100.4^{\circ}\text{F}/38^{\circ}\text{C}$ , respiratory distress, cyanosis, persistent vomiting, lethargy, dehydration, worsening jaundice; maternal heavy bleeding, severe headache, chest pain/SOB, wound infection, suicidal ideation.
- **Acceptance:** Scenario tests route correctly 100% of the time; “call 911/ED now” style messaging appears when triggered.
- **Physician-vetted Knowledge Base (KB)** with ingestion & versioning of PDFs, DOCX, HTML (including NayaCare blog/pages); required metadata:
  - `topic`, `audience` (infant|birthing parent), `age-window` (e.g., 0-7d, 8-28d, 1-3mo), `risk-level` (general|high-risk), `summary`, `author`, `reviewed-by`, `review-date`, `expiry-date`, `source-url`, `tags`.
  - **Acceptance:** Upload → parse → searchable within 10 minutes; version history retained; stale items auto-flagged on `expiry-date`.
  - **RAG pipeline** (retrieval-augmented generation) that strictly prefers KB content; fallback to general references only for low-risk, generic education and always cited as such.
  - **Acceptance:** At least 80% of answers cite KB items when available; hallucination rate <2% in review set.
  - **Persona & tone:** empathetic, calming, plain-language ( $\leq$ 8th-grade reading level), culturally sensitive; names Dr. Patel as supervising author of materials when applicable.
  - **Acceptance:** Automated readability check  $\leq$  Grade 8 for 95% of messages; tone guidelines unit tested on 30 canned prompts.
  - **PDF generator** producing personalized handouts + visit summary with: timestamp, topic outline, vetted content excerpts, links to KB, “what’s normal vs. red flags,” self-care checklist, follow-up instructions.
  - **Acceptance:** PDF renders in <5s; accessible tagging (PDF/UA where feasible); downloadable link expires in  $\leq$ 24h by default.
  - **HIPAA posture (MVP):** data minimization (no PHI unless user explicitly opts-in), TLS 1.2+ in transit, AES-256 at rest, RBAC, audit trail, configurable retention (default 0-7 days), BAAs with all vendors, breach notification workflow.
  - **Acceptance:** Security review checklist passes; BAAs executed; audit logs evidence access and admin actions.
  - **Admin console** for Dr. Patel/staff: content upload, tagging, approval workflow, version diff, quick “safety notice” banner push, analytics dashboard (de-identified), and conversation redaction tools.
  - **Acceptance:** Role-based access; editor → reviewer → publish states enforced.
  - **Observability & QA:** conversation transcripts (de-identified by default), flagged-content queue, scenario-test harness ( $\geq$ 50 curated neonatal & maternal scenarios), latency/error dashboards.
  - **Acceptance:** P95 response  $\leq$ 8s; uptime  $\geq$ 99.5% monthly; safety tests run on each release.
  - **Legal & clinical disclaimers** contextualized by location (practice state), coverage hours, and emergency options; explicit consent statement when collecting any PHI.
  - **Acceptance:** Dynamic banner reflects state and hours; consent stored with timestamp & policy version.

## Should Have (S)

- **Age-aware guidance** based on baby’s DOL/gestational age and maternal postpartum day; asks minimal questions to personalize.
- **Spanish (es-US) content parity** for top 20 topics; locale toggle persisted.

- **Voice input** (Web Speech API) with on-device transcription; opt-in only.
- **Email/SMS secure delivery** of PDFs via expiring links; patient verification via magic-link.
- **Soft handoff to on-call staff** (message dropbox) with SLA banner; no live chat required for MVP.
- **Topic maps** (guided flows) for common concerns: feeding amounts, latch pain, spit-up vs. vomiting, normal stools, safe sleep, jaundice, postpartum mood, incision care.
- **Basic analytics**: top intents, unresolved questions, KB coverage gaps, satisfaction CSAT prompt.

## Could Have (C)

- **Photo intake** (e.g., rash, umbilical stump) with strict consent + PHI notice; runs only on-device classification or vendor with BAA; never stored by default.
- **EHR integration** (read-only) for patient age and care plan summaries via FHIR; behind auth wall.
- **Proactive nudges**: day-by-day tips for first 28 days (opt-in).
- **Multi-model ensemble** (Claude + GPT + Perplexity orchestrated) with cost/latency gating and safety voting.

## Won't Have (W)—for MVP

- Diagnosis, prescriptions, dosing advice beyond Dr. Patel-approved dosing charts.
- Open community chat or user-generated medical content.
- Unbounded data retention.

## Non-Functional Requirements

- **Performance**: P95 end-to-first-token  $\leq$ 3s; P95 full-answer  $\leq$ 8s.
- **Security**: OWASP ASVS L2; SAST/DAST in CI; secret rotation; IP allowlist for admin.
- **Reliability**: 99.5% uptime/month; graceful degradation if LLM vendor down; multi-region storage for KB.
- **Compliance**: HIPAA administrative/physical/technical safeguards documented; BAAs in place; quarterly access reviews.
- **Content quality**: Every KB item has `review-date  $\leq$  12 months`; lint checks for readability & inclusive language.

## Acceptance Test Matrix (excerpt)

- 50 red-flag scenarios → 100% correct escalation messaging.
- 100 common questions →  $\geq$ 95% helpfulness (clinician-rated),  $\leq$ 2% hallucination.
- 10 PDF samples → all pass accessibility checks; cite correct versioned KB entries.