PulseMed Business Framework

White-Label Healthcare Technology Platform for Medical Practices

Your Practice's AI Integrator — Built by Physicians, For Physicians

Document Type: Framework
Date: December 2025
Classification: Internal Strategy & Enablement
Audience: Business Partners (Franchisees) & non-executives (Referral Sources for healthcare environment)

Executive Summary

PulseMed provides white-label, physician-controlled healthcare technology platforms to medical practices. Our platform delivers ego-protected, support-focused care at an **exponential cost-reduction** (~1/100th of traditional solutions) — a language unknown in healthcare, or physician-approved content. This solves the critical access problem that prevents patients from receiving care they need but cannot afford.

The Problem We Solve

Commercial AI in healthcare "Large Mess, no clinical ROI" — Generic tools lack clinical context, creating:

- Poor scope-specific clinical outcomes (high hallucination rates in ER)
- Data sovereignty violations (PHI leakage)
- No accountability (no liability)
- No HIPAA-compliance guarantees
- No audit trail or accountability

Result: Physicians don't trust AI, patients get unreliable information, and practices can't leverage AI to reduce workload.

Our Solution: Three-Pillar Architecture

- **Physician-Curated Knowledge Base:** AI uses ONLY content uploaded and approved by the practice's own physicians
- **Ego-Protected Support:** Patients receive empathetic, non-judgmental support, reducing burden on clinical staff
- **Sustainable Innovation:** AI tools are tested by overburdened, ego-protected staff before deployment to patients

Optional Features (commercial per practice):

- **Algorithmic Triage:** Rules-based or clinical decision support for appropriate care recommendations (for practices that want it)

Technical Differentiators

- **Production-grade technology:** Not theoretical; already deployed in live healthcare environments
- **Data sovereignty:** PHI never leaves practice control
- **Scope-integrated accountability:** Tools aligned with organizational liability
- **Physician-controlled:** Practices decide what AI knows and how it responds
- **HIPAA-compliant:** Built for healthcare from the ground up
- **GDPR-compliant:** Ready for international expansion
- **Continuous Business Development:** Ongoing integration (internal to each practice) ensures specific and specific needs
- **Integrated admin dashboard:** Easy to use AI prompting, video tutorials, feedback mechanisms, data collection, and real-time quality monitoring
- **Sustainable innovation:** Single vendor supporting hundreds of practices, not thousands

Market Opportunity

- **Target:** Mixed practices, health centers, specialty clinics
- **Pain Point:** Access barriers under private pay/concierge models
- **Value:** Reduce after-hours calls (for PA), staff burnout, patient retention/satisfaction

Problem Statement

Why Commercial AI Fails in Healthcare

1. The Information Overload Problem

Commercial AI tools lack the critical context:

- Poor environmental/periodontal (operative)
- Medical/drug-drug and herbal (Lexicomp)
- Outdated information (dangerous)
- Hallucinations (liability)

This AI must "decide" which scenario to read data. This creates:

- A risk infrastructure where no standard exists
- Inconsistent responses to identical questions
- No way to guarantee evidence-based answers

2. The Control Problem

Physicians have zero control over:

- What information the AI uses
- How the AI phrases a response
- When the AI escalates to a human
- Which clinical guidelines it follows

3. The Liability Problem

- No accountability framework
- No audit trail for decisions
- No HIPAA-compliance
- Legal exposure unknown

The PulseMed Solution

We eliminate this problem through **curated architecture**:

Problem	Commercial AI	PulseMed
Knowledge Source	Entire internet	Physician-approved documents only
Hallucination Rate	6.4% - 14%	<1%
Physician Control	None	Complete
Audit Trail	None	Every response traceable
HIPAA Compliance	No	Built-in

Solution Overview

What We Provide

A complete, white-label healthcare technology platform that practices deploy to their patients.

1. Knowledge Base Management

- Physicians upload approved content (PDFs, guidelines, protocols)
  - Automatically text extraction and chunking
- AI ONLY references this knowledge base (no internet)
- Content organization and version control

2. AI Engine

- Google FHIR-approved knowledge base connector
- Patient-friendly translation (Flesch-Kincaid 6th-8th grade)
- Scope-specific responses (PA, NP, Dentist)
- Multi-modal support (Gemini 2.0Pro, gpt-4o)

3. Clinical Triage System (Optional)

- **Availability for practices that want it:** Not all clinics require triage functionality
- Rules-based or AI-assisted
- Activity scoring algorithm
- Escalation protocols (yes / no / urgent-but-safe / emergency / ultra-urgent)
- Clinical decision support integration
- **Example use case:** Pediatric practices avoid triage for urgent symptoms
- **Example use case:** Orthopedic practices may want aggressive triage (bunion surgery)

4. Patient Interface

- White-label chatbot (practice branding)
  - Mobile-responsive design
- Sustainable onboarding for retirees
- Mobile-only patient access

5. Integrated Admin Dashboard

- **Easy to use interface:** Designed for non-technical staff
- **Prompt engineering without coding:**
  - Visual scenarios (business logic/decision trees, drag-and-drop)
  - Pilot scenarios are predefined (disabled)
  - Control branching and messaging
  - Test in sandbox before deployment
- **Data collection control:** Practices decide:
  - What patient data is collected
  - Conversation logging
  - PHI/de-identification
  - Privacy and consent settings
- **Real-time monitoring:** Live conversation feed
  - Escalation alerts
  - Performance metrics
  - Response analytics (including
  - Usage
  - Conversation analytics
- **Knowledge base management:** Upload, edit, delete, and organize documents
- **Conversation analytics:** Review patient interactions and AI responses
- **Quality control:** Ensure AI provides accurate, practice-approved information
- **Survey and feedback tracking:** Patient satisfaction and future usage analytics

6. Data & Analytics

- Patient engagement metrics
- Response quality tracking
- Usage frequency
- ROI reporting

Current Features & Integrations

- **Specialty-specific workflows:** Tailored features for different medical specialties (pediatrics, cardiology, oncology, etc.)
- **Multi-modal AI:** Supports Google Gemini 2.0Pro, gpt-4o, and other models as needed
- **Compact integrations:** Focused AI products specific tools, HIPAA-compliant, designed for specific workflows
- **Continuous development:** Build specialized features that accommodate specific physician needs and needs
- **Flexible architecture:** Platform designed to extend and adapt to new practice needs

How It Works

```
Patient asks question

AI retrieves relevant physician-approved documents

Generates patient-friendly response using ONLY approved content

Applies triage algorithm if enabled for this practice

Logs conversation for quality assurance
```

Note: Practices have (e.g., colleagues) response with concerns because they are the platform punch for (ensure navigation without things functionality).

The Answer "walled garden": Access response comes from physician-approved sources.

The Three-Pillar Architecture

Pillar 1: Curated Knowledge Base

What: A physician-curated repository of medical information

Why: Prevents hallucinations

How It Works:

- Physicians upload their documentation (PDFs, protocols, handouts)
- Systems extracts text (OCR for scanned files)
- Content is chunk/indexed, cleaned for retrieval
- AI queries ONLY from this approved corpus

Technical Implementation:

- Document storage (Google Cloud Bucket or equivalent)
- Text extraction (pdf-parse / Tesseract OCR)
- Chunking strategy (500-1000 tokens, paragraph boundaries)
- Relevance scoring (keyword + recency, TBD: semantic embeddings)
- Top-3 most relevant documents sent to AI per query

Why It Works:

- Eliminates hallucinations (no unverified sources)
- Practices provide specific priorities
- Self-audit tool (every response traceable)
- Physicians maintain control

Pillar 2: Progressive Learning

What: The system improves from interactions with continuous quality

How It Works:

- Patient conversations are logged (with audit trail)
- Physicians review AI responses periodically
- Gaps (knowledge is assumed or absent) are identified
- Approved content is integrated into knowledge base
- Quality metrics are continuously monitored

Quality Controls:

- Physicians approve all new content
- Hallucinations are tracking (<1% target)
- Source attribution is mandatory (no unsourced claims)
- Patient satisfaction is monitored

Why It Works:

- Practices control what AI learns
- Quality improves over time (fewer gaps)
- Institutional knowledge is captured
- Scalable (more practices = more data = better AI)

Pillar 3: Humanistic Interaction

What: A compassionate, multilingual, ego-protected patient interface

How It Works:

- Conversational tone optimized for healthcare contexts
- Emotionally intelligent responses (for stressed patient needs
- Clear, scope-specific language (no technical jargon)
- Context-aware (empathetic acknowledgments, referrals, escalations)
- Multilingual capability (serves diverse patient populations
- Ego-protected (avoids critical judgments with past poor experiences

Technical Implementation:

- Response prompts emphasize empathy and warmth
- Multi-turn conversation memory (maintains context across messages)
- Escalation triggers (detects when patient needs human support)
- Fallback responses (when AI uncertain, defaults to "contact your clinic")

Why It Matters:

- Patients feel heard and understood
- Reduces anxiety by avoiding overwhelming medical jargon
- Builds trust through consistent, caring communication
- Improves patient satisfaction and engagement
- Differentiates from competitors (clinical focus only)

Proven Technical Framework

All components are production-tested, deployed in healthcare environments, and designed for physicians. These are not theoretical capabilities—they are **proven innovations** that differentiate PulseMed from competitors.

1. HIPAA-Compliant De-Identification Layer

All patient queries are automatically scrubbed of personally identifiable information (PII) **before** being sent to any AI provider:

- SSN detection and masking: `/\b\d{3}[-\s]?\d{2}[-\s]?\d{4}\b/g`
- Phone number masking: `/\b\d{10}\b/g`
- Email address masking: `/\b[A-Z0-9._%+-]+@[A-Z0-9.-]+\.[A-Z]{2,}\b/gi`
- Date masking: `/\b(0[1-9]|1[0-2])[-/](0[1-9]|[12]\d|3[01])[-/](19|20)\d{2}\b/g`

Why It Matters: Even if using third-party AI providers, no PHI ever leaves the practice's control. This is automatic, requiring zero configuration from the practice.

2. Intelligent Query Classification System

Not every patient question needs the same response approach. Our system automatically classifies queries:

- **Medical question** → Retrieve relevant knowledge base documents
- **Emotional support question** → Prioritize empathy and reassurance over clinical information
- **Emergency question** → Trigger immediate escalation protocols

Practice Customization: Practices can define custom patterns (pediatric urgent keywords, cardiac-specific terms, etc.)

Why It Matters: Patients receive the right type of response—clinical when they need facts, emotional when they need reassurance, without requiring practitioners to manually triage every message.

3. Emergency Detection & Escalation Protocol

Multi-layered emergency detection ensures patient safety:

```
Emergency keywords detected:
→ SMS + Immediate E-Link notification (response time: 30 min-1hr)
→ AI continues normal conversation
```

Smart Limitations: Prevents false positives (e.g., "pain pills" doesn't trigger an "injury")

Why It Matters: AI handles routine questions 24/7, but physicians are immediately notified when patients need human attention.

4. Intelligent Document Retrieval with Relevance Scoring

When a patient asks a question, the system doesn't just dump all documents to the AI. It uses a **relevance scoring algorithm**:

Match Type	Score
Exact match	+100 points
Partial match	+50 points
Keyword match	+10 points

Top-3 most relevant documents are sent to the AI, sorted by score.

Multi-source Unified Retrieval:

- Primary storage (Firebase File Storage)
- Backup knowledge base
- Unified search

Why It Matters: The AI gets exactly the right context for each question, not a random selection. Better accuracy, lower token costs.

5. Temporal Awareness (Conversation Memory)

The AI remembers conversations for 30 minutes (configurable):

- Prevents repeating identical responses
- Recognizes follow-up questions ("what about my other medication?")
- Continuously tracks conversation flow ("this is the fourth time we've talked about this")

Why It Matters: Conversations feel natural, patients don't have to repeat themselves, and the AI builds on what was already discussed.

6. Clinical Photo Analysis Framework

For practices that want visual assessment capacity, we've a strong clinical photo analysis:

- **Expert-crafted prompts:** Lesion morphology, color, texture, distribution
- **Structured diagnostic reasoning:** Different differentials for different patient presentations
- **Confidence calibration:** Honest about uncertainty, recommends specialist when unclear
- **Red flag recognition:** Emergency vs. urgent vs. routine
- **Patient-friendly responses:** Translates clinical findings into plain language

Why It Matters: Extends the platform's capabilities beyond text into visual assessment. Practices can offer visual assessment guidance.

7. Hybrid PDF Processing (Text + OCR)

Many practices have scanned documents that standard text extraction can't read. Our hybrid approach:

- Attempts fast text extraction (free, for text-based PDFs)
- Falls back OCR processing (for scanned/image-based documents)
- Supports:
  - Medical charts
  - Handwritten notes
  - Lab reports
  - Insurance or legal documents

Why It Matters: Practices don't need to know whether their PDFs are text-based or scanned. Upload anything, it just works.

8. Multi-Profile Patient Management

One practice can track family accounts or complex patient relationships:

- Multiple profiles per account
- Profile-specific log filtering
- Each linked account (parents, caregivers, family members)
- Profile switching with session isolation

Why It Matters: Pediatric groups, caregiver support, complex families and patients relationship dynamics.

9. Comprehensive Tracking & Pattern Recognition

Built for logging systems with automatic insight generation:

- **Activity logging:** Visit types, interventions, and behaviors
- **Pattern recognition:** Identify trends (peak usage times, common questions)
- **Quality metrics:** Response accuracy, escalation rates, satisfaction
- **Exported reports:** Analysis-ready data (for quality improvement)

Why It Matters: Practices gain actionable insights from patient interactions, not just raw data.

10. Embeddable Widget Architecture

Deploy the chatbot anywhere with a simple script tag:

```
<script src="https://yourclinic.pulsemed.ai/embed.js"></script>
```

Features:

- Shadow isolated for security
- Mobile-responsive design
- Customizable branding/colors
- Secure conversation isolation
- Works on any website, patient portal, or EHR with interface

Why It Matters: Practices can deploy to their existing website, patient portal, or EHR integration.

11. Consent Versioning & Compliance Tracking

Full legal compliance tracking with version control:

- Legal documents version tracking (TOS, Privacy Policy, Consent)
- Automatic re-consent prompts when versions change
- Consent logging with timestamps
- Compliance audit trail (who consented to what, when)

Why It Matters: When legal documents are updated, every affected patient is automatically re-prompted. Full audit trail for compliance reviews.

12. Research-Ready Analytics

Data structure prepared for quality improvement or research:

- **Feature usage tracking:** Which features are used most
- **Conversation data collection:** With consent, research-quality data
- **Sentiment analysis:** Positive/negative/neutral classification
- **Topic extraction:** Automatic categorization of patient concerns

Why It Matters: Practices can demonstrate outcomes, participate in research, and continuously improve based on data.

LLM Comparison & Scaling

Model Options

Model	Medical Accuracy	Speed	Cost (1K conversations)	Best For
Gemini 2.0Pro	★★★★★	★★★★	$0.50-$1.00	Medical accuracy
GPT-4 Turbo	★★★★★	★★★	$1.00-$2.00	Complex reasoning
GPT-4o	★★★★★	★★★★★	$0.25-$0.50	Balanced
Claude Haiku	★★★★	★★★★★	$0.08-$0.15	Speed / volume
GPT-4.1-mini	★★★★	★★★★	$0.05-$0.10	Budget-friendly

Recommended Strategy

Phase 1 (Launch): Gemini Pro 1.5

- Proven medical accuracy
- Strong HIPAA-compliant infrastructure
- Multi-modal capabilities

Phase 2 (Scaling): Hybrid Routing

- Complex queries/triage → Gemini Pro
- Simple queries → GPT-4.1-mini or Haiku
- Cost savings: 60-70%

Phase 3 (Enterprise): Multi-Model Architecture

- Emergency/triage → Gemini Pro (best accuracy)
- Medical questions → Claude or GPT-4
- General questions → Haiku
- Customer-specific preferences

Cost Optimization

- **Query classification:** Route simple queries to cheaper models
- **Response caching:** Cache common answers
- **Context optimization:** Minimize token usage
- **Volume discounts:** Negotiate pricing at scale

Marketing & Positioning

Core Message

"Stop Lady of Hospital Healthcare Platform That Physicians Actually Trust"

Key Messaging

DO Say:

- "Physician-owned" / "physician-approved"
- "Your practice's knowledge, AI powered"
- "Ego-protected patient support, not judgment"
- "Zero hallucinations risk"
- "Author after Hours calls by the doctor"
- "HIPAA-compliant from day one"

DON'T Say:

- "Replace your staff" / "replace PA/ nursing"
- "AI that replaces doctors" (you supporter)
- "Autonomous medical decisions" (liability)
- "Revolutionary AI" (buzzy, skeptic)

Value Propositions by Audience

First Physicians:

- Reduce after-hours calls by 60%
- Protect consistent patient education
- Ego-protected patient support
- Full audit trail for liability protection
- **Continuous business development:** Ongoing integration ensures specific needs and usage of identified
- **Integrated admin dashboard:** Easy to use, video tutorials, feedback mechanisms, data collection, and real-time monitoring for quality control

First Practice Managers:

- **Full accountability tracking (integrated admin dashboard):** Easy to use UI with complete conversation logs, feedback, and real-time monitoring
- Improve patient satisfaction scores
- Reduce administrative burden
- Measurable ROI
- Scalable across locations

Front Desk Staff:

- Less stress in patient questions
- Clear escalation protocols
- Reduce repetitive questions
- Transition assistants when needed
- Focus on what

Competitive Positioning

Competitor	Their Focus	Our Advantage
Generic AI chatbots	Broad knowledge	Physician-curated, healthcare-specific
Telehealth platforms	Replace visits	Augment existing practice
Scribe tools	Documentation	Patient-facing, not provider-facing
Chatbot vendors	Lead generation	Clinical focus, ego-protected

Business Model

Revenue Streams

- **Subscription:** Monthly platform access
- **Implementation:** Initial setup and training
- **Customization:** Custom workflows, branding
- **Integration:** EHR connectors, API access

Pricing Tiers

- **Basic:** $[illegible]/month — Single provider, 500 conversations
- **Professional:** $[illegible]/month — Multi-provider, 2,000 conversations, priority support
- **Enterprise:** Custom — Unlimited, dedicated support, SLA

Unit Economics

- **Cost per conversation:** $0.[illegible] (Gemini API)
- **Gross margin:** [illegible]%
- **Customer acquisition cost:** [illegible]
- **Lifetime value:** [illegible]

Implementation Roadmap

Phase 1: Foundation (Weeks 1-4)

- Knowledge base upload
- Prompt configuration
- Widget deployment
- Staff training

Phase 2: Optimization (Weeks 5-8)

- Monitor conversations
- Identify gaps
- Refine responses
- Measure impact

Phase 3: Scale (Ongoing)

- Add new knowledge
- Expand use cases
- Multi-location rollout
- Continuous improvement

Risk Mitigation

Risk	Mitigation
Hallucinations	Strict prompts, source attribution, confidence scoring
Data breach	No PHI storage, encryption, audit logs
Regulatory change	Flexible architecture, compliance monitoring
Competition	Physician ownership, clinical focus, relationships
Adoption resistance	Staff training, gradual rollout, champions

Contact & Next Steps

Demo: [illegible]
Pilot: [illegible]
Partnership: [illegible]

vs. ChatGPT/Bard: "Unlike commercial AI, we use only your approved content

vs. EHR Chatbots: "More accurate, more customizable, better experience, full autonomy via integrated admin dashboard"

Business Model

Pricing Tiers

Starter - $2,500/month

- Up to 500 conversations/month
- 100 documents in knowledge base
- **Triage protocols** (optional - can be disabled)
- Email support
- Single location

Professional - $5,000/month

- Up to 2,000 conversations/month
- 500 documents
- **Custom triage protocols** (optional - can be disabled)
- Priority support
- Up to 3 locations
- Basic analytics

Enterprise - $12,000/month

- Unlimited conversations
- Unlimited documents
- Fully custom protocols
- **Custom feature development** - Specialty-specific modules and unique integrations
- **Practice-specific workflows** - Tailored to physician's scope of medicine
- Dedicated support
- Unlimited locations
- Advanced analytics
- EHR integration
- SLA guarantees

Implementation Fees

Standard Setup: $5,000 one-time

- Knowledge base configuration
- **Triage protocol setup** (if needed - optional)
- Branding customization
- 2-hour training
- Go-live support

Custom Setup: $10,000-25,000

- Complex EHR integration
- **Custom feature development** - Specialty-specific modules, unique integrations, practice-specific workflows
- **Unique integrations** - Connect with practice-specific tools, scheduling systems, or third-party services
- Multi-location setup

Unit Economics

Per Practice (Professional tier):

- Revenue: $5,000/month
- LLM costs: ~$1,500/month (variable)
- Infrastructure: ~$300/month
- Support: ~$400/month
- **Gross margin: ~56%**

At scale (200+ practices), margins improve to 65-70%.

Revenue Projections

Year	Practices	Monthly Revenue	Annual Revenue
1	20	$80,000	$960,000
2	75	$300,000	$3,600,000
3	200	$800,000	$9,600,000

Implementation Roadmap

> ***Note:*** *The core platform is already built and production-proven. This roadmap focuses on packaging, client onboarding, and scaling.*

What's Already Built (Production-Ready)

- ✅ AI chat engine with physician-curated knowledge base
- ✅ Admin dashboard with full resource management
- ✅ PDF processing (text extraction + OCR fallback)
- ✅ HIPAA-compliant de-identification layer
- ✅ Emergency detection and escalation protocols
- ✅ Intelligent relevance scoring algorithm
- ✅ Embeddable widget (single script tag deployment)
- ✅ Consent versioning and compliance tracking
- ✅ Analytics and research-ready data collection
- ✅ Multi-profile patient management
- ✅ Temporal awareness (conversation memory)
- ✅ Clinical photo analysis framework

Phase 1: Platform Generalization

- Configure multi-tenancy (separate knowledge bases per practice)
- Build configuration-driven branding (logos, colors, practice names)
- Create standardized deployment pipeline
- Document onboarding process

Phase 2: Client Onboarding (Per Practice)

Each new client follows this process:

1. **Discovery call** - Understand practice specialty and needs
2. **Feature configuration** - Enable/disable modules (triage, photo analysis, etc.)
3. **Knowledge base setup** - Upload practice's documents and protocols
4. **Branding customization** - Logo, colors, chatbot name
5. **Embed code deployment** - Single script tag on practice website
6. **Training** - 2-hour staff orientation
7. **Go-live support** - First week monitoring

Phase 3: Scale

- Standardize sales process
- Build marketing website and case studies
- Develop partnership channels (EHR vendors, medical associations)
- Expand to 20+ practice deployments

Custom Development (Only When Requested)

- EHR integrations specific to practice
- Specialty-specific features
- Unique workflow automations

Competitive Analysis

Why We Win

Competitor Type	Their Weakness	Our Advantage
Commercial AI (ChatGPT)	Not healthcare-specific, hallucinations	Curated knowledge, physician control
EHR Chatbots	Generic, poor UX	Practice-specific, modern AI, custom features
Telehealth Platforms	Not 24/7, expensive	Always available, complements care
Medical AI Startups	No physician control, one-size-fits-all	Full curation, control, and custom development

Our Moat

1. **Production-proven technology** - Not vaporware; 12 core innovations already battle-tested in live healthcare environments
2. **First-mover with curated approach** - Nobody else does this at scale
3. **Physician relationships** - Trust built through control and transparency
4. **Custom feature development** - Unique integrations and specialty-specific modules create practice-specific value
5. **Integrated admin dashboards** - Easy-to-use UI giving clinics full autonomy creates deep operational dependency
6. **Integration depth** - Hard to switch once embedded in workflows with custom features
7. **Progressive learning** - System gets smarter, competitors can't replicate history
8. **Automatic HIPAA compliance** - PHI de-identification happens before data leaves practice control

---

Risk Assessment

Technical Risks

Risk	Impact	Mitigation
LLM provider outage	High	Multi-model support, caching
Hallucination despite curation	High	Quality monitoring, source attribution
Scalability	Medium	Cloud infrastructure, horizontal scaling

Business Risks

Risk	Impact	Mitigation
Slow adoption	High	Free pilots, strong case studies
Regulatory changes	Medium	HIPAA compliance built-in, legal counsel
Competition	Medium	First-mover advantage, deep integration
LLM cost increases	Medium	Multi-model routing, caching

Legal Risks

Risk	Impact	Mitigation
Liability for AI responses	High	Disclaimers, source attribution, escalation
HIPAA violations	High	Compliance from day one, regular audits

---

Summary

Why PulseMed Wins

1. **Production-proven technology** - 12 core innovations already deployed in live healthcare environments
2. **Solved the accuracy problem** - Curated knowledge eliminates hallucinations
3. **Automatic HIPAA compliance** - PHI de-identification before data leaves practice control
4. **Physician control** - Practices trust what they control
5. **Humanistic interaction** - Emotionally appropriate AI that doesn't overwhelm patients
6. **Custom feature development** - Unique integrations and specialty-specific modules tailored to each practice's scope of medicine
7. **Integrated admin dashboards** - Easy-to-use UI giving clinics full autonomy over resources, data collection, and real-time quality monitoring
8. **Scalable architecture** - One framework, unlimited practices
9. **Clear ROI** - 60-80% after-hours call reduction
10. **Deploy anywhere** - Single script tag embeds chatbot on any website or patient portal

The Opportunity

Every medical practice wants AI that:

- Reduces workload
- Provides accurate information
- Maintains physician control
- Integrates with existing workflows
- **Adapts to their specific specialty and practice needs** - Custom features and unique integrations
- **Gives them full autonomy** - Easy-to-use admin dashboard for complete control over resources, data collection, and real-time quality monitoring

PulseMed delivers all of this.

The three-pillar architecture (curated knowledge + progressive learning + humanistic interaction) is our competitive advantage. Combined with optional features like algorithmic triage (when needed), this flexible approach allows us to serve diverse practice needs—from simple resource navigation AI usage to full-featured clinical decision support systems that are specifically crafted for the medical scope.

Document End

Proprietary & Confidential - PulseMed LLC